

LEARN CODING

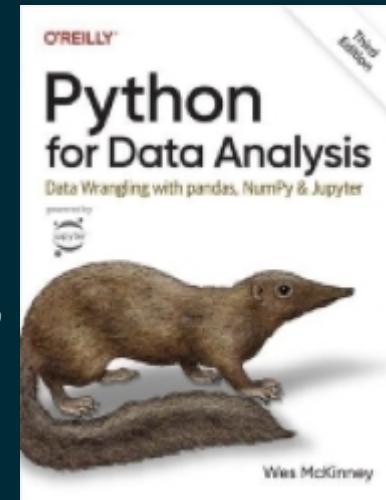
ale66

PANDAS

- Created by Wes McKinney, a ‘quant’ for hedge-fund AQR.
- a library for processing tabular data, both numeric and time series.
- it provides data structures (series, dataframe) and methods for data analysis.

```
1 pip install pandas
```

W. McKinney, Python
for Data Analysis, 3/e.
O'Reilly 2022.



DATA STRUCTURES - SERIES

SERIES

A one-dimensional object containing values and associated labels, called Index.

Unless we assign indices, Pandas will simply enumerate the items.

```
1 import numpy as np  
2 import pandas as pd  
  
1 # a simple series  
2 s1 = pd.Series([10, 20, 30, 40])  
3  
4 s1
```

```
0    10  
1    20  
2    30  
3    40  
dtype: int64
```

```
1 # Assign explicit indices to our data
2 s2 = pd.Series([10, 20, 30, 40], index = ['a', 'b', 'c', 'd'])
3
4 s2
```

```
a    10
b    20
c    30
d    40
dtype: int64
```

Example: putting 10 quid a month into a savings account

```
1 my_savings = pd.Series([10, 20, 30, 40],
2                         index = ['jan', 'feb', 'mar', 'apr'])
3
4 my_savings
```

```
jan    10
feb    20
mar    30
apr    40
dtype: int64
```

From dictionaries to Pandas series

```
1 # keys correspond to indices.  
2 my_dict = {'a':10, 'b':20, 'c':30, 'd':40}  
3  
4 s3 = pd.Series(my_dict)  
5  
6 s3
```

```
a    10  
b    20  
c    30  
d    40  
dtype: int64
```

Use the index to select one or more specific values.

```
1 # Get the data on position 'a' of s3  
2  
3 s3['a']
```

```
np.int64(10)
```

```
1 # Get the data indexed 'a' and 'c' of s3  
2  
3 s3[['a', 'c']]
```

```
a    10  
c    30  
dtype: int64
```

Filter elements

```
1 # Select data which is Less than 25  
2  
3 s3[s3<25]
```

```
a    10  
b    20  
dtype: int64
```

apply element-wise mathematical operations...

```
1 # Square every element of s3  
2  
3 s3**2
```

```
a    100  
b    400  
c    900  
d   1600  
dtype: int64
```

or a combination of both:

```
1 # Square every element of s3 smaller than 25  
2  
3 s3[s3<25]**2
```

```
a    100  
b    400  
dtype: int64
```

DATA STRUCTURES - DATAFRAMES

DATAFRAMES

2D structures where values are labelled by their index and column location.

```
1 # Notice how we specify columns.  
2 new_df = pd.DataFrame([10, 20, 30, 40],  
3                         columns = ['Integers'],  
4                         index = ['a', 'b', 'c', 'd'])  
5 new_df
```

Integers

a 10

b 20

c 30

d 40

```
1 # Implicitly add a column.  
2 new_df['Floats'] = (1.5, 2.5, 3.5, 4.5)  
3  
4 new_df
```

	Integers	Floats
a	10	1.5
b	20	2.5
c	30	3.5
d	40	4.5

DATA STRUCTURES: DATAFRAME loc

Select data according to their location label.

```
1 # here Loc slices data using index name.  
2  
3 new_df.loc['c']
```

```
Integers    30.0  
Floats      3.5  
Name: c, dtype: float64
```

```
1 # here Loc slices data using column name.  
2  
3 new_df.loc[:, 'Integers']
```

```
a    10  
b    20  
c    30  
d    40  
Name: Integers, dtype: int64
```

alternatively, use `new_df['Integers']`

Loc queries can be combined:

```
1 # here we use both index and column name.  
2  
3 new_df.loc['c', 'Integers']
```

```
np.int64(30)
```

DATA STRUCTURES: DATAFRAME - iloc

Select a specific slice of data according to its position (index).

```
1 # here loc slices data using index number.  
2 new_df.iloc[2]
```

```
Integers    30.0  
Floats       3.5  
Name: c, dtype: float64
```

```
1 # here loc slices data using column number.  
2 new_df.iloc[:, 0]
```

```
a    10  
b    20  
c    30  
d    40  
Name: Integers, dtype: int64
```

```
1 # here we use both index and column number.  
2 new_df.iloc[2, 0]
```

```
np.int64(30)
```

DATA STRUCTURES: DATAFRAME - FILTERS

Complex selection is achieved applying Boolean filters.
Multiple conditions can be combined in one statement.

```
1 new_df[new_df['Integers'] > 10]
```

	Integers	Floats
b	20	2.5
c	30	3.5
d	40	4.5

```
1 # here we apply conditions to both columns.  
2  
3 new_df[(new_df.Integers>10) & (new_df.Floats>2.5)]
```

	Integers	Floats
c	30	3.5
d	40	4.5

DATA STRUCTURES: DATAFRAME - Axis

DataFrames operate on 2 dimensions.

Axis = 0 invokes functions across rows

default behaviour when the axis is not specified.

```
1 new_df.sum()
```

```
Integers    100.0
Floats       12.0
dtype: float64
```

Axis = 1 invokes functions across columns.

```
1 new_df.sum(axis=1)
```

```
a    11.5
b    22.5
c    33.5
d    44.5
dtype: float64
```

We can mix element-wise operations with axis functions

Example: Create a column with the sum of squares of each row.

```
1 # Just one line of code!
2 new_df['Sumsq'] = (new_df**2).sum(axis=1)
3 new_df
```

	Integers	Floats	Sumsq
a	10	1.5	102.25
b	20	2.5	406.25
c	30	3.5	912.25
d	40	4.5	1620.25

FROM NUMPY TO PANDAS

FROM NUMPY TO PANDAS: `where()`

In Numpy, the `where()` allows to describe actions associated to `True` and `False`
an if/then/else construct, essentially

```
1 l = np.arange(9).reshape((3, 3))  
2 l
```

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

```
1 # If True then make it double, else halve it  
2 np.where(l<5, l*2, l/2)
```

```
array([[0. , 2. , 4. ],  
       [6. , 8. , 2.5],  
       [3. , 3.5, 4. ]])
```

P. executes `where()` differently: when `False` it assigns n/a

```
1 df_1 = pd.DataFrame(1)
2 df_1
```

	0	1	2
0	0	1	2
1	3	4	5
2	6	7	8

```
1 df_1.where(df_1<5)
```

	0	1	2
0	0.0	1.0	2.0
1	3.0	4.0	NaN

NUMPY FUNC. TO PANDAS OBJECTS

```
1 # l is a Numpy matrix which readily interoperates with Pandas
2 my_df = pd.DataFrame(l, columns=['A', 'B', 'C'])
3
4 my_df
```

	A	B	C
0	0	1	2
1	3	4	5
2	6	7	8

```
1 # Extract the square root of each el. of column B (NB: my_df remains unchanged)
2 np.sqrt(my_df.B)
```

```
0    1.000000
1    2.000000
2    2.645751
Name: B, dtype: float64
```

BACK AND FORTH B/W PANDAS AND NUMPY

```
1 # Extract the values back into a Numpy object  
2 m = my_df.values  
3 m
```

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

IMPORTING DATA

Read a datafile and turn it into a DataFrame. Several arguments are available to specify the behavior of the process:

`index_col` sets the column of the csv file to be used as index of the DataFrame

`sep` specifies the separator in the source file

`parse_dates` sets the cols. to be converted into *datetimes*

```
1 FILE = './path/to/some_file.csv'  
2  
3 df_r = pd.read_csv(FILE,  
4                     index_col = 0,  
5                     sep = ';',  
6                     parse_dates = ['date'] )
```

BIOSTATS DATA - info()

The `info()` method outputs top-down information on the DataFrame

```
1 MYDATA = 'data/biostats.csv'  
2  
3 df_bio = pd.read_csv(MYDATA)  
4  
5 df_bio.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 18 entries, 0 to 17  
Data columns (total 5 columns):  
 #   Column      Non-Null Count  Dtype     
---  --          --          --       --  
 0   Name        18 non-null    object    
 1   Sex         18 non-null    object    
 2   Age         18 non-null    int64    
 3   Height(in)  18 non-null    int64    
 4   Weight(lbs) 18 non-null    int64  
dtypes: int64(3), object(2)  
memory usage: 852.0+ bytes
```

BIOSTATS DATA - head() AND tail()

Handy visualisation of first/last n rows (default = 5)

```
1 df_bio.head()
```

	Name	Sex	Age	Height(in)	Weight(lbs)
0	Alex	M	41	74	170
1	Bert	M	42	68	166
2	Dave	M	32	70	155
3	Dave	M	39	72	167
4	Elly	F	30	66	124

```
1 df_bio.tail()
```

	Name	Sex	Age	Height(in)	Weight(lbs)
13	Neil	M	36	75	160
14	Omar	M	38	70	145
15	Page	F	31	67	135
16	Luke	M	29	71	176
17	Ruth	F	28	65	131

BIOSTATS DATA - INDEX COLUMN

Selecting the index column affects the structure of the DataFrame and thus information retrieval.

Caution: the index does not have to be unique. Multiple rows could have the same index name.

```
1 # here we set the Name column as the index
2 df_bio2 = pd.read_csv(MYDATA, index_col = 0)
3
4 df_bio2.head(2)
```

Name	Sex	Age	Height(in)	Weight(lbs)
Alex	M	41	74	170
Bert	M	42	68	166

```
1 #It is now possible to use elements of the Name column to select an entire row
2 df_bio2.loc['Bert']
```

```
Sex          M
Age         42
Height(in)   68
Weight(lbs)  166
Name: Bert, dtype: object
```

DESCRIPTIVE STATS - `describe()`

Compute the descriptive stats of quantitative variables

```
1 # Descriptive statistics for the Age variable  
2 df_bio['Age'].describe()
```

```
count    18.000000  
mean     34.666667  
std      7.577055  
min      23.000000  
25%     30.000000  
50%     32.500000  
75%     38.750000  
max      53.000000  
Name: Age, dtype: float64
```

`try df_bio.describe()`

DESCRIPTIVE STATS - CATEGORICAL VARS

The `value_counts()` method computes the unique values and how many times they occur.

```
1 # Descriptive statistics for the entire DataFrame  
2 df_bio.Sex.value_counts()
```

```
Sex  
M    11  
F     7  
Name: count, dtype: int64
```

PANDAS DATA DISPLAYS

Pandas objects come with methods for visualisation they are built on top of matplotlib

```
1 df_bio['Age'].plot(kind = 'hist')  
2  
3 # alternative syntax:  
4 # df_bio.Age.plot(kind = 'hist')
```

